

Robert Schumacher

Microsoft

Don't package your libraries,
write packageable libraries!




Foreword

- This talk will be presented entirely from a packaging and maintainer viewpoint
- Every point presented has other tradeoffs that need to be evaluated for your project
- My objective is that you will be better engineers by being aware of the packaging implications of your design decisions

Foreword

- This talk will include real code from real projects
- This is not intended to shame – All of the projects mentioned are awesome and make all our lives easier.
- A handful of choices that make *packaging* harder can still be the right ones for the project to make.
- I sincerely thank all authors for their endless hours of effort to make these projects possible and hope they continue into the future!

Four Facets of the Packaging Ecosystem





Maintainers

- These are the volunteers which will package your library for you
- They are not contributors to your project
- They do not even use your project
- They do not know your dev workflows
- They do not know your code
- They are human and will make mistakes



Buildsystems

Use a popular buildsystem

- Today, that's CMake, MSBuild, or Autoconf
- As a maintainer, I can handle your buildsystem being bad in the same ways that everyone else is bad.
 - I *can't* handle your buildsystem being bad in its own new and clever ways.

The reward for using a buildsystem before it's mainstream is to get packaged poorly or not at all.

Use a popular buildsystem (cont.)

- Boost: Boost.Build
- OpenSSL: 5.3k+ lines of Perl
 - https://github.com/openssl/openssl/blob/OpenSSL_1_0_2p/Configure
- FFmpeg: 7.4k+ lines of Bash
 - <https://github.com/FFmpeg/FFmpeg/blob/efb65abedf40c0a5bc6eb76e6cf19b633a143444/configure>
- Qt: Qmake

*You discover that one of the above doesn't build if there are spaces in the path.
Good Luck!*

Header-only is not a Packaging Panacea

- "What's hard about copying a few headers around?"
 - How will you generate config.cmake/.pc files?
 - How will I run your tests?
- You lose the ability to encapsulate
- You are forcing "static linkage" for consumers
- It's extremely easy to end up with circular dependencies between header-only libraries
- In the wild:
 - Boost v1.68 has *multiple* circular dependencies between header-only components
 - <https://pdimov.github.io/boostdep-report/>
 - Shout out to pdimov for his work on analyzing the boost internal dependencies!

Use your buildsystem's standard constructs

- CMake:
 - ``BUILD_SHARED_LIBS``
 - ``find_package()``
 - ``find_dependency()``
 - Use targets instead of macros!
- In the wild:
 - Libpng v1.6.35 uses ``PNG_SHARED`` and ``PNG_STATIC``
 - <https://github.com/glennrp/libpng/blob/v1.6.35/CMakeLists.txt#L70>
 - Expat v2.2.6 uses ``BUILD_shared``
 - https://github.com/libexpat/libexpat/blob/R_2_2_6/expat/CMakeLists.txt#L34

Do not use header-checks

- This is highly susceptible to contamination
 - Example:
 - You see `openssl.h` and assume OpenSSL 1.1 is available
 - In fact, it was provided by libressl
 - You try to use some openssl-specific thing and fail to build
 - I didn't even want you to build with SSL anyway! ☹
- Partial support becomes poisonous
 - Nobody can provide a partial `unistd.h` or `sys/mmap.h` because your build will suddenly believe it's on Linux instead of Windows and fail


Do not use header-checks (cont.)

- Libxml2 v2.9.4 checks for ``dirent.h``, ``finite()``, ``fpclass()``, ``fp_class()``, ``isnan()`` then does nothing with them
 - <https://github.com/GNOME/libxml2/blob/bdec2183f34b37ee89ae1d330c6ad2bb4d76605f/configure.ac#L561>
 - Fixed in latest master 🙌
- Exiv2 2018-09-18 checks for ``mmap()`` and, if found, uses it on Windows and fails
 - <https://github.com/Exiv2/exiv2/blob/f06b69fa56051204ef3b8f1379ed2a0346672d22/src/basicio.cpp#L436>

Do not “detect” different settings


- Give the builder options to require that a setting is ON or OFF
- Respect those settings – don’t undermine them
 - CMake note: every use of `find_package()` without `REQUIRED` is suspect
- In the wild:
 - Cpprestsdk v2.10.6 uses an old, embedded copy of websocketpp if `find_package()` quietly fails
 - https://github.com/Microsoft/cpprestsdk/blob/v2.10.6/Release/cmake/cpprest_find_websocketpp.cmake#L6-L13

Remember: The maintainer will make mistakes! Don’t leave spikes in pits 😊



Don't depend
on having a
host compiler

- Most common offender: building code generators during the build
- If you must depend on the host, have separate build stages that can be packaged independently and imported
- In the wild:
 - See previous custom buildsystems – several are built during their parent builds
 - Protobuf v3.6.1 does this perfectly with `--with-protoc=` ✨
 - <https://github.com/protocolbuffers/protobuf/blob/v3.6.1/src/README.md>



Don't hardcode a single package manager

- This includes the `ThirdParty/` “package manager” 😊
- All package managers need to intercept your dependencies to handle ODR, updates, security, consistency, and more
- Hardcoding any particular package manager makes this *much* more difficult

Do not use runtime plugins

- They always imply custom build and deployment logic *in all consuming projects* that is extremely difficult and fragile
- In the wild:
 - Qt5 (when dynamic linking) has many plugin folders which need to be handled with extreme special casing (see: `windeployqt`)
 - <https://github.com/qt/qttools/blob/5.11/src/windeployqt/main.cpp#L799-L828>

Don't muck with compiler flags

- CMake: `CMAKE_CXX_FLAGS`
- Add an escape hatch if you must
- In the wild:
 - Caffe2 v0.8.1 forces `/MT` instead of `/MD` when building static libs on MSVC
 - <https://github.com/caffe2/caffe2/blob/v0.8.1/CMakeLists.txt#L96-L105>
 - Fixed in master! 🐛

Dependencies





Don't have them

Hide dependencies



Don't include ``Windows.h`` in your public headers

Too much macro configuration that is impossible to encapsulate (WIN32_LEAN_AND_MEAN, WINSOCK, etc)



In the wild:

OpenSSL v1.0.x pulled in ``windows.h`` inside ``rand.h``

- https://github.com/openssl/openssl/blob/OpenSSL_1_0_2p/crypto/rand/rand.h
- Fixed in v1.1 🐞

Do not have optional dependencies

- Provide optional features as separate libraries with separate builds
- The user doesn't always know what they want; rebuilding everything every time they change their mind is awful

Do not have optional dependencies (cont)



OpenCV v3.4.1 optionally depends
on pretty much every graphics
format library written before 2013

<https://github.com/Microsoft/vcpkg/blob/7881abfc29c916330e868118b29606cb32c51b16/ports/opencv/CONTROL#L1-L78>



FFmpeg says “Hold my beer”

<https://github.com/FFmpeg/FFmpeg/blob/release/4.0/configure#L203>

Adopt newer dependencies quickly

- If you use Boost, you should be on 1.68 right now
- If your dependency releases once a year
 - You take 6mo to adopt it
 - Someone else takes 3mo to adopt you
 - Someone else 1.5mo to adopt that
 - Someone else ...
- Users can use N-1 just barely before N is released.

We must move faster.

ABI, API & Support



Don't replace APIs when using C++17 vs 14

- For all major compilers and STLs, newer standards are additive and compatible
- We do this to make it easy for users to upgrade their code incrementally
- This means that you need one binary that can serve *both* versions

Note: only applies if you support both C++14 and C++17.

Don't replace APIs when using C++17 vs 14

- Most commonly, this is broken by optional polyfills.
- Example:
 - Abseil provides a custom `absl::string_view` or typedefs `std::string_view` based on the c++ standards level
 - https://github.com/abseil/abseil-cpp/blob/e01d95528ea2137a4a27a88d1f57c6cb260aafed/absl/strings/string_view.h#L33-L41
 - <https://github.com/abseil/abseil-cpp/blob/2a62fbdedf64673f7c858bc6487bd15bcd2ca180/absl/base/config.h#L354-L365>

Don't worry about ABI; worry about API

- Avoid breaking source compatibility and provide migration guidance if you do
 - Ideally in a way that can be applied by someone who isn't familiar with either codebase
- If you must maintain ABI compatibility
 - Your public binary interface should be C
 - You should use caller-allocated memory
 - You should not expose any dependencies in your interface

Have a clear, single supported source version

- If you must support multiple versions, make them simultaneously consumable
 - Separate headers, separate symbols
- Multiple supported versions make it harder for others to depend on you
 - If X, Y, and Z are “supported” for customers to be using today, which should I test with? Will every downstream library make that same conclusion?

Have a clear, single supported source version

- Qt has multiple supported versions in flight
 - And has the macros to (theoretically) avoid ODR violations!
 - <https://github.com/qt/qtbase/blob/c5307203f5c0b0e588cc93e70764c090dd4c2ce0/src/corelib/global/qglobal.h#L164-L188>
- OpenSSL is simultaneously supporting v1.1.x and v1.0.x
 - Changed the library names but not the header names

Use a standard license

- When you use a standard license, it is infinitely easier for everyone downstream to know exactly what to do
- Make sure that license is readily accessible in a dedicated file
- In the wild:
 - Sqlite3 is in the "public domain" (Surprisingly, not well defined)
 - No dedicated license file!
 - <https://sqlite.org/copyright.html>

Code



Break source, not behavior

- v1: ``void throws_on_error()``
- v2: ``std::error_code throws_on_error() noexcept``
- Users: 😞
- Nobody tracks test code coverage *of their dependencies*

Namespace *all the things*

- Symbols with ``namespace``
- Headers with ``foo/``
- In the wild:
 - Libjpeg-turbo v1.5.3 exports ``/jconfig.h``
 - Nlohmann-json v3.0.1 exported ``/json.h``
 - <https://github.com/nlohmann/json/blob/v3.0.1/README.md>
 - Fixed in v3.1.0 🍰

Don't be hostile to the compiler

- ``#if _MSC_VER > 1900` + `#error``
- ``#if !$COMPILER` + `#error``
- ``#if _MSC_VER > 1900` + `#message SPAM``
- In the wild:
 - PMDK v1.4.2 hardcoded a failure for ``_MSC_VER` > 1911`
 - <https://github.com/pmem/pmdk/blob/1.4.2/src/common/util.h#L287-L291>
 - Fixed in master! (at least, lifted to `>=2000`) 🚀

Patterns






Patterns

- Maintainers will make mistakes
- Packaging is more than your library + your user
- Many effects are hard to see locally

Shameless advertisement: vcpkg!

- All examples in this talk were taken from vcpkg's 750+ catalog of libraries for Mac, Windows, and Linux
- Our community deals with these problems so users don't need to 😊
- Help us out by sending a PR!
- <https://github.com/Microsoft/vcpkg>

Take our survey <https://aka.ms/cppcon>



You can win an Xbox One S - Starter Bundle



9/27 10:30 – 12:00 // Breckenridge Hall

Thoughts on a More Powerful and Simpler C++ (5 of N), Herb Sutter

Other sessions

Monday, September 24th

- ~~14:00 – 15:00~~
~~How to Write Well Behaved Value Wrappers~~
~~by Simon Brand~~
- ~~15:15 – 16:15~~
~~How C++ Debuggers Work~~
~~by Simon Brand~~

Tuesday, September 25th

- ~~14:00 – 15:00~~
~~What Could Possibly Go Wrong?: A Tale of Expectations and Exceptions~~
~~by Simon Brand and Phil Nash~~
- ~~15:15 – 15:45~~
~~Overloading: The Bane of All Higher Order Functions~~
~~by Simon Brand~~

Wednesday, September 26th

- ~~12:30 – 13:30~~
~~C++ Community Building Birds of a Feather~~
~~with Stephan T. Lavavej and others~~
- ~~14:00 – 15:00~~
~~Latest and Greatest in the Visual Studio Family for C++ Developers 2018~~
~~by Marian Luparu and Steve Carroll~~
- ~~15:15 – 15:45~~
~~Don't Package Your Libraries, Write Packagable Libraries!~~
~~by Robert Schumacher~~

Wednesday, September 26th

- ~~15:15 – 15:45~~
~~What's new in Visual Studio Code for C++ Development~~
~~by Rong Lu~~
- 15:50 – 16:20
Value Semantics: Fast, Safe, and Correct by Default
 - by Nicole Mazzuca
- 16:45 – 17:45
Memory Latency Troubles You? Nano-coroutines to the Rescue! (Using Coroutines TS, of Course)
 - by Gor Nishanov
- 18:45 – 20:00
Cross-Platform C++ Development is Challenging – let Tools Help!
 - by Marc Goodner and Will Buik

Thursday, September 27th

- 9:00 – 10:00
Inside Visual C++'s Parallel Algorithms
 - by Billy O'Neal
- 15:15 – 15:45
ConcurrencyCheck – Static Analyzer for Concurrency Issues in Modern C++
 - by Anna Gringauze
- 16:45 – 17:45
Class Template Argument Deduction for Everyone
 - by Stephan T. Lavavej

Don't assume your user's link model

- Don't use ``LoadLibrary`/`dlopen``
- Don't use ``DllMain``
- Glib on Windows uses ``DllMain()`` which prevents static linking :'
 - <https://github.com/GNOME/glib/blob/2.58.1/glib/glib-init.c#L273-L294>
- Don't use mutable globals

Default to portability

- If you support SSE2 and AVX, default to SSE2
- This is because portability is *extremely* hard to test
 - The maintainer will not see your flag because everything seems to run fine
 - The user who's hunting for every perf gain possible *will* find your flag
- In the wild:
 - Rocksdb v5.14 defaults to ``-march=native` / `/arch:AVX2``
 - <https://github.com/facebook/rocksdb/blob/5.14.fb/CMakeLists.txt#L194-L210>

Don't make me build things I don't need

- Samples, docs, alternate flavors, tools
- A crowd favorite is to combine this with dep autodetection
 - "If I can find doxygen, then generate my docs by default"
 - This translates to "The build runs 200% slower on that machine that appears identical"
- In the wild:
 - Fast-RTPS v1.6.0 builds and installs examples with no easy OFF
 - <https://github.com/eProsima/Fast-RTPS/blob/v1.6.0/CMakeLists.txt>

Don't build with ``-Werror`/`/WX`` by default

- Your user will use a newer compiler. Your code will break. Your user will be sad. 😞
- Note: Definitely turn it on for your CI and dev work!
- In the wild:
 - Folly 2018.09.24 throws ``/Werror`` unconditionally on unix
 - <https://github.com/facebook/folly/blob/v2018.09.24.00/CMake/FollyCompilerUnix.cmake#L19>

Golden rule(s) of source layout:

- Lay out your sources like you don't have a buildsystem
- Lay out your sources like you have 10 buildsystems

Ship often

- It is the package manager's job to help the user manage change
- Shipping more gives users more choice: half now, half later, instead of everything later.
- Minimizing delta between “last stable” and “master” enables:
 - Backporting changes
 - Useful bug reports
 - Fewer fixed-in-master-but-unavailable bugs

Have clear public headers