

Mini Dumps

Efficient core dumps for FlashBlade



Outline

- Motivation
- Simple fixes that don't fix it (MADV_DONTDUMP)
- Core file, signal handlers
- Summary of solution
- The devil is in the details
- Caveats and future work

Motivation

- Core dumps are super useful for debugging
- But kernel-driven core dumps dump everything in memory:
 - Customer data, secret keys, etc.
 - 30+ seconds to dump
 - 54GB dumps (before compression)
- Desire for faster dumps
- Desire for smaller dumps (with little loss of debuggability)
- Desire for eliminating regions from the dump
- But it still needs to be readable using gdb

Why not MADV_DONTDUMP?

- Not a safe function call at dump time
 - So we would pay a cost at runtime for an event that may never happen
 - Kernel lock to protect data structure
- Interactions with memory allocator
 - We allocate 2MB pages in 64MB chunks and re-use them
 - No-dump regions change quickly
- Applies at at least 4k granularity
 - We may want finer-grained control since important data and customer data could share a page
- Still many GB dumped we probably don't need

What's in a core file?

- `readelf -a core` shows us the contents

ELF header

```
e_ident = MAGIC,  
    64-bit, LSB  
e_type = ET_CORE  
e_machine = x86_64  
program hdr size  
    phdr count  
section hdr cnt = 0
```

Program headers

```
1 PT_NOTE  
    info on NT_foo ent  
N PT_LOAD  
    (memory entries)  
    mode, vaddr, size  
    p_filesz (can be 0)  
    P_offset (can be 0)  
    (0 means valid addr  
    but no contents)
```

“Notes”

```
NT_PRPSINFO (uid, gid, pid)  
NT_SIGINFO (signo, errno)  
NT_AUXV (auxv_t from OS)  
NT_FILE: count, page_size, []  
    start, end addr, offset  
    (/proc/self/maps)  
Per-thread:  
    NT_PRSTATUS (signo, tid,  
    regs)  
NT_FPREGSET, NT_X86_XSTATE
```

Signal handlers

- Program errors: `SIGSEGV`, `SIGINT`, `SIGILL`, possibly others
- `std::abort: SIGABRT`
- We can catch all of these
- But can we do enough in a signal handler?
 - `man 7 signal`
 - Safe function calls: `open()`, `read()`, `write()`, `mkdir()`, etc.
 - Most file I/O can be done in a signal handler
 - No memory allocation though
- We can do file I/O, and we know what a `core` file looks like

Summary: What to dump?

- Need Per-thread state
 - Registers at dump time
 - Signal info for crashing thread
- “Interesting” memory contents
 - Stack contents (how far back?)
 - Memory near any register value that looks like a pointer
 - Memory near any stack value that looks like a pointer
 - Also chase anything in those blocks that may be a pointer?
 - How much around each “anchor” address?

Summary: What to dump?

- Main thread info comes from signal handler
- Send signal to other threads to dump their per-thread info
- One thread must dump the memory contents and memory map

So that's really it -- we have a signal handler which pokes all the other threads (SIGUSR1); we read `/proc/self/maps`, and we write out interesting memory contents

Details: Opt-in

- Signal handler only registered via `ps_init_platform()` call to `setup_signal_handlers()`
 - Need copy of `argv` for `NT_PRPSINFO`'s `pr_psargs` field
- Need to reserve memory for unknown quantities
 - `PT_LOAD` info for dump contents, `NT_FILE` info for memory map
 - Need one per line in `/proc/self/maps`, plus splitting when not fully dumped
 - Assume we have at most 128GB memory and map in 2MB chunks; 64k entries max
 - Space for per-thread info
 - $2 * k_max_execution_count$
 - Bounce buffer for memory contents that can't be written directly: 4kB
 - String storage for file names from `/proc/self/maps` since we read that file in chunks

Details: `signal_received()`

1. Check for cascading signals and exit if reentrant too many times
2. `fflush()` stdout, stderr; `setvbuf` to non-buffered
 - Not documented as signal safe but seems to work
3. Signal other threads (only on first entrance to handler)
4. Dump registers, stack to stdout for the logs
 - Need `sig_printf()` [format to local buffer then `write(STDOUT_FILENO)`] to be signal safe
 - Internally we also tell minidump about stack range, registers
5. One thread dumps `/proc/self/maps` to logs
6. Wait for all threads to tell minidump about their registers + stack
7. Write minidump file
8. Wait for all threads before re-signalling or returning so we get correct exit status

Details: signalling other threads

- Need to know all threads that are running, so we can signal at dump time
 - Since we start threads, we can memoize `pthread_t` and OS thread id (`gettid()`) on start
 - `pthread_kill()` other threads
 - `SIGUSR1` is unused for other purposes
1. Print registers to log
 2. Tell minidump about our `siginfo_t`
 3. Print stack to log
 4. Print some stack contents to log
 5. Determine valid stack range and tell minidump
 6. Spin forever, `sleep(100)`

Details: per-thread valid stack range

- Limit to 64kB
 - We just need some reasonable number
 - Allocated stack is 8MB; sigaltstack is 512kB
- We could memoize valid stack addresses on thread start
- Currently we “probe” by writing to a nullfd
 - Use /dev/random instead of /dev/null so the handler isn't a no-op
 - write() returns ≥ 0 , it was a valid address

Details: `/proc/self/maps`

- We parse each line of `/proc/self/maps`
 - Read 1kB at a time (need some fixed buffer size)
 - Print to log on crash
 - Pass to minidump to parse text contents
- Each line looks like
 - `<hex begin>-<hex end> <permissions> <hex offset> <dev-id> <inode> <filename>`
 - Filename is optional, could be special tags like `[stack]`
 - Permissions are `rwX[ps]` with `-` for no permission; `'p'` rivate or `'s'` hared
 - Offsets indicate different mapped portions of a file
 - Ignore devid and inode

Details: `/proc/self/maps`

- Memoize the range and name, determine initial dump plan
 - Don't save filename for `/anon_hugepage`
 - Determine “real” file based on leading `/` (e.g. `/lib/x86_64-linux-gnu/libc-2.1.19.so`)
 - Never dump unreadable or shared mappings (we don't use shared mmap)
 - Always dump files
 - in theory file content can be determined from symbols, but it's only a few 100kB
 - Always dump special regions, e.g. `[stack]`, `[vdso]`, `[vsyscall]`
 - `[vsyscall]` requires `memcpy` then `write(2)`
 - Assume entries not 2MB aligned are ‘interesting’ and always dump
 - I don't recall what's here but it's not too many bytes
 - Default no-dump everything else
 - DMA memory, malloc contents
 - We will only dump things from anchor addresses

Details: anchor addresses

- From a value that looks like a memory address, ensure data from [addr - before, addr + after) is in the dump
 - `add_anchor_addr()` for each register, adding 1k before and 8k or 1M after
 - `add_anchor_addr()` for each word in the stack, adding 128 bytes before and 1k or 4k after
 - Larger ranges for the faulting thread
- Keep all known memory mappings in sorted order; look for entry overlapping anchor address (4k aligned)
 - We don't use convenient containers like `std::map` because they allocate on insert
 - We could provide an allocator that draws from a pre-allocated pool of nodes
 - Split entry if needed
 - Check if next entry is adjacent and shift 4k from one to the next (common case)
 - Be careful with `offset` field for files

Details: write minidump file in order

- core file has an ordering for each section
 - We emit per thread and per process info in the same order
 - Crashing thread first
- Log how much data we kept from anchor addresses, how long minidump took

INFO: Keeping 59344 kB from mapped files (e.g. shared libs) and 259292 kB from other sources; **elided 1232792 kB**

INFO: Adding anchor addresses for 15 threads in 252 maps

INFO: **Added 1796 kB** from anchor addresses

INFO: Cannot open directory /ssd/core (errno 2); falling back to "."

INFO: Writing minidump for 15 threads, 456 mappings

INFO: Emit NT_SIGINFO for LWP 21449

INFO: Emit NT_FILE for 130 entries, 9398 bytes

INFO: minidump took **1639 ms**

Details: process monitoring

- Process may appear up even if heartbeat is non-responsive
 - It takes non-zero time to write a minidump
 - `kill -9` and restart during this time would be tragic
- Give a 10 second extra buffer for dump to succeed before kill + restart

Caveats

- We currently have no sparse virtual addresses
 - So if it's in `/proc/self/maps`, and it's one of our 2MB chunks, it's also physical
 - This is a solvable problem since we own the memory allocator

Future work

- API to mark ranges as “never dump” or “always dump”
 - Desire to only mark things “always dump” at startup, not forward operation
 - Per-server or per-authority data structures that are of high value
 - And to not mark too much as “always dump”
 - Desire to rarely mark regions as “never dump”
 - AES keys no longer in process
 - Mark DMA buffers and messaging buffers as no-dump?
 - Lock-free crash-safe table? We can’t afford to segfault when iterating
- Save `NT_FPREGSET` and `NT_X86_XSTATE` regions too?
 - So far these haven’t been missed

Summary

- Signal handler can write a valid `core` file
- Signal other threads to get their registers and stack at crash time

From an early crash in a `runtestsD` test:

```
$ ls -l minidump core
-rw----- 1 mdf mdf 1.4G Jun 18 18:22 core
-rw-r--r-- 1 mdf mdf 313M Jun 18 18:22 minidump
```

- Usable core dumps in an order-of-magnitude less space and time