

---

---

# What do you *mean* "thread-safe"?

— Geoff Romer —

---

---

# What does "thread-safe" mean?

"A thread-safe function can be safely invoked concurrently with other calls to the same function, or with calls to any other thread-safe functions, by multiple threads."

— POSIX

# What does "thread-safe" mean?

```
int in[100], out[100];
```

```
void thread1() {  
    ...  
    memcpy(&out, &in,  
           sizeof(in));  
    ...  
}
```

```
void thread2() {  
    ...  
    memcpy(&out, &in,  
           sizeof(in));  
    ...  
}
```

# What does "thread-safe" mean?

```
int in[100];
```

```
void thread1() {  
    ...  
    int out[100];  
    memcpy(&out, &in,  
           sizeof(in));  
    ...  
}
```

```
void thread2() {  
    ...  
    int out[100];  
    memcpy(&out, &in,  
           sizeof(in));  
    ...  
}
```

# What does "thread-safe" mean?

The Talk in One Slide

const

== mutable

**== thread safe**

(bitwise const or  
internally synchronized)



— [Herb Sutter](#)

# What does "thread-safe" mean?

The Talk in One Slide

const  
== mutable

wat

**== thread safe**  
(bitwise const or  
internally synchronized)



— [Herb Sutter](#)

**What are we trying to be "safe" from?**

# What are we trying to be "safe" from?

Race conditions?

"A race condition or race hazard is the behavior of an electronics, software, or other system where the **output is dependent on the sequence or timing** of other uncontrollable events." — [Wikipedia](#)



# What are we trying to be "safe" from?

Data races?

```
int i = 0;
```

```
void thread1() {  
    ..  
    ++i;  
    ..  
}
```

```
void thread2() {  
    ..  
    std::cout << i;  
    ..  
}
```

# What are we trying to be "safe" from?

Data races?

```
std::string s = "";
```

```
void thread1() {  
    ...  
    s.append("foo");  
    ...  
}
```

```
void thread2() {  
    ...  
    std::cout << s;  
    ...  
}
```

# What are we trying to be "safe" from?

## Data races?

"The execution of a program contains a *data race* if it contains two potentially concurrent conflicting actions, at least one of which is not atomic...

Two expression evaluations *conflict* if one of them modifies a memory location (4.4) and the other one reads or modifies the same memory location."

— [The C++ Standard](#)

# What are we trying to be "safe" from?

Data races?

```
std::string s = "";
```

```
void thread1() {  
    ...  
    s.append("foo");  
    ...  
}
```

```
void thread2() {  
    ...  
    std::cout << s;  
    ...  
}
```

# What are we trying to be "safe" from?

## API races

An *API race* occurs when the program performs two concurrent operations on the same object, when the object's API contract doesn't permit those operations to be concurrent.

# Identifying API races

```
Widget shared_widget;
```

```
void thread1() {  
    ...  
    shared_widget.foo();  
    ...  
}
```

```
void thread2() {  
    ...  
    shared_widget.bar();  
    ...  
}
```

# Identifying API races

```
Widget shared_widget;
```

```
void thread1() {  
    Thingy t;  
    ...  
    t.foo(shared_widget);  
    ...  
}
```

```
void thread2() {  
    Whatever w;  
    ...  
    w.bar(shared_widget);  
    ...  
}
```

# Identifying API races

If a live object has a **thread-safe type**, it can't be the site of an API race.



# Identifying API races

```
// Thread-safe
class JobRunner {
    JobSet running_;
    JobSet done_;
    std::mutex m_;

    void OnJobDone(Job* job) {
        m_.lock();
        running_.erase(job);
        done_.insert(job);
        m_.unlock();
    }
};
```

```
// Thread-safe
class JobSet {
    std::set<Job*> jobs_;
    std::mutex m_;

    void erase(Job* job) {
        m_.lock();
        jobs_.erase(job);
        m_.unlock();
    }
};
```

# Identifying API races

```
int shared_int;
```

```
void thread1() {  
    Thingy t;  
    ...  
    t.foo(shared_int);  
    ...  
}
```

```
void Thingy::foo(int i);
```

```
void thread2() {  
    Whatever w;  
    ...  
    w.bar(shared_int);  
    ...  
}
```

```
void Whatever::bar(  
    const int& i);
```

# Identifying API races

```
Widget shared_widget;
```

```
void thread1() {  
    Thingy t;  
    ...  
    t.foo(shared_widget);  
    ...  
}
```

```
void Thingy::foo(  
    const Widget& widget)
```

```
void thread2() {  
    Whatever w;  
    ...  
    w.bar(shared_widget);  
    ...  
}
```

```
void Whatever::bar(  
    const Widget& widget);
```

# Identifying API races

If a live object has a **thread-safe type**, it can't be the site of an API race.

If a live object has a **thread-compatible type**, it can't be the site of an API race if it's not being mutated.

# Identifying API races

```
class LazyStringView {  
    const char* data_;  
    mutable optional<size_t> size_;  
    mutable std::mutex mu_;  
  
public:  
    size_t size() const {  
        std::scoped_lock lock(mu_);  
        if (!size_)  
            size_ = strlen(data_);  
        return *size_;  
    }  
};
```

# Identifying API races

```
Widget shared_widget;
```

```
void thread1() {  
    Thingy t;  
    ...  
    t.foo(shared_widget);  
    ...  
}
```

```
void thread2() {  
    Whatever w;  
    ...  
    w.bar(shared_widget);  
    ...  
}
```

# Identifying API races

If a live object has a **thread-safe type**, it can't be the site of an API race.

If a live object has a **thread-compatible type**, it can't be the site of an API race if it's not being mutated.

If a live object has **any type**, it can't be the site of an API race if it's not being accessed concurrently.

# Identifying API races

```
struct Counter {  
    int c = 0;  
    void operator()() { ++c; }  
};  
const std::function<void()> f = Counter{};  
  
void thread1() {  
    f();  
}  
  
void thread2() {  
    f();  
}
```



# Identifying API races

```
Widget shared_widget;
```

```
void thread1() {  
    Thingy t;  
    ...  
    t.foo(shared_widget);  
    ...  
}
```

```
void thread2() {  
    Whatever w;  
    ...  
    w.bar(shared_widget);  
    ...  
}
```

# Identifying API races

```
void Thingy::foo(  
    const Widget&) {  
    ...  
    baz();  
    ...  
}  
  
void baz() {  
    static int counter = 0;  
    counter++;  
}
```

```
void Whatever::bar(  
    const Widget&) {  
    ...  
    baz();  
    ...  
}
```

# Identifying API races

If a live object has a **thread-safe type**, it can't be the site of an API race.

If a live object has a **thread-compatible type**, it can't be the site of an API race if it's not being mutated.

If a live object has **any type**, it can't be the site of an API race if it's not being accessed concurrently.

... but beware of **thread-hostile functions**, which can cause API races at sites other than their inputs.

# Identifying API races

A given line of code is guaranteed to have no API races if it calls no thread-hostile functions, all inputs are live, and each input is

- not being accessed by other threads, or
- thread-safe, or
- thread-compatible and not being mutated by any thread.

# Identifying API races

```
vector<int> shared_vec = {0, 0};
```

```
void thread1() {  
    ...  
    ++shared_vec[0];  
    ...  
}
```

```
void thread2() {  
    ...  
    ++shared_vec[1];  
    ...  
}
```

# Identifying API races

```
vector<bool> shared_vec = {false, false};
```

```
void thread1() {  
    ...  
    shared_vec[0] = true;  
    ...  
}
```

```
void thread2() {  
    ...  
    shared_vec[1] = true;  
    ...  
}
```

# Identifying API races

```
vector<int> shared_vec = {0, 0};
```

```
void thread1() {  
    ...  
    ++shared_vec[0];  
    ...  
}
```

```
void thread2() {  
    ...  
    ++shared_vec[1];  
    ...  
}
```

# Identifying API races

```
// Increments every value in the range [begin, end).
template <typename Iterator>
void f(Iterator begin, Iterator end) {
    for (Iterator it = begin; it != end; ++it)
        ++*it;
}
```

```
std::vector<int> v = {1, 2, 3};
```

```
void thread1() {
    f(v.begin(),
      v.begin() + 2);
}
```

```
void thread2() {
    f(v.begin() + 1,
      v.end());
}
```



# Identifying API races

```
class Widget {  
    int* counter_;  
  
public:  
    Widget(int* counter)  
        : counter_(counter) {}  
  
    // Thread-hostile!  
    void Twiddle() {  
        ++*counter_;  
    }  
};
```

```
Widget MakeWidget();  
  
void thread1() {  
    Widget w = MakeWidget();  
    w.Twiddle();  
}  
  
void thread2() {  
    Widget w = MakeWidget();  
    w.Twiddle();  
}
```

# Recommendations

For library code:

- Make your types thread-compatible if possible, thread-safe if necessary.
- Clearly document any types that are thread-safe, or thread-incompatible.
  - Prefer to explicitly document the rest as thread-compatible.
- Be thoughtful about directly exposing sub-objects.
- Never define or use thread-hostile functions.
  - Avoid hidden mutable shared state
  - Be very careful with private pointers to shared data.

For application code:

- Make shared objects thread-safe, or thread-compatible and immutable.

# Questions?